

Notes on "The BrainScaleS-2 Accelerated Neuromorphic System With Hybrid Plasticity"

Piyush Sud

8/22/2024

1 Abstract

- This paper describes the BrainScaleS SNN, which has both analog and digital components.

2 Introduction

- BrainScaleS provides continuous time modeling of the brain using a physical replica of neurons connected with synapses.
- BrainScaleS uses CMOS instead of nano-devices for this continuous modeling to create configurable behavior.
- This system addresses the time and temperature drift problem of most analog systems.
- On chip calibration is achieved with embedded processors and ADCs.
- This chip can be operated in many different modes - the first is batch mode, where there are no data dependencies between experiments.
- This chip supports fine control of synaptic plasticity.
- Analog data is acquired in parallel, while the plasticity changes are computed efficiently using a digital processor.
- Vector matrix multiplication can be realized for ANNs.

3 The BrainScaleS-2 System

System Architecture

- The system designed so far can be considered as a single core, which can theoretically be used as a building block for a large SNN.
- A single core consists of:
 - 4 256x128 synaptic crossbars, each containing 128 neurons, with one neuron corresponding to a single column.
 - Two digital processors for controlling the plasticity of the synapses.

- Two 512 channel ADCs. Each quadrant has 256 channels, since 128 channels are used for spike train correlation measurements and the other 128 channels are used for spike train anti-correlation measurements.
 - Analog parameter storage
 - An event routing network for spike communication. The outputs of the neurons can either be routed back to the inputs or to an external output.
- This chip has been used in two hardware setups, one for analog measurements and the other for edge devices, with a Zynq FPGA SoC.
 - In the future, an external event routing network could be used with 8 cores on a single chip.

Accelerated Analog Emulation of Neural Dynamics

- The speed of the semiconductors is around 1000x the speed of biological neural networks in order to work with the time constants of semiconductor materials.
- The neurons are implemented using mixed-signal circuits (it's unspecified what the exact circuit is), which follow the adaptive exponential integrate-and-fire (ADEx) model.
- The first equation says that $C \frac{dV}{dt}$ = the leakage current into the extracellular fluid + the current caused by positive feedback of Na⁺ ions - the adaptation/refractory current after a spike + the current I flowing into the membrane through external stimuli, where C is the capacitance of the membrane and V is the membrane potential. E_l is the resting potential of the extracellular fluid.

$$C_m \dot{V} = -g_l(V - E_l) + g_l \Delta_T \exp\left(\frac{V - V_T}{\Delta_T}\right) - w + I, \quad (1)$$

$$\tau_w \dot{w} = a(V - E_L) - w, \quad (2)$$

•

- The second equation describes the adaptation current after a spike has been fired, which decays to 0 with time constant τ_w .
- Each neuron can be configured with 80 bits of SRAM and 24 analog parameters which are set by a 10-bit DAC. Since each neuron can be controlled individually, production deviations can be compensated for to create an ideal network. Parameters such as the refractory period and membrane capacitance can be programmed.
- Each neuron integrates the current from all 256 synapses.
- The following equation describes the current flowing into a neuron, and is equal to the weighted sum of the spike trains multiplied by an exponential decay.

$$I_{\text{syn}} = \sum_i w_i S_i(t) * e^{-t/\tau_{\text{syn}}}$$

-
- The weights are stored locally per synapse in a 6-bit SRAM.
- Instead of using muxes to route synapses, a 6-bit address is associated with each synapse. If there is a connection between the output of a neuron and a synapse, then any "event packets" that are sent on the output of that neuron will contain the address of the synapse, which will check for a matching address and respond only if the address matches.
- Analog sensor circuits on each synapse can be used to accurately measure the spike timing, which in turn can be used for STDP.
- By disabling the spiking behavior, a regular non-time continuous VMM can be performed.

Hybrid Plasticity and Versatile Digital Control

- The PPU (Plasticity Processing Unit) is designed to be very flexible to allow for many different digital configurations.
- The ASIC can perform tens of thousands of correlation measurements per second.

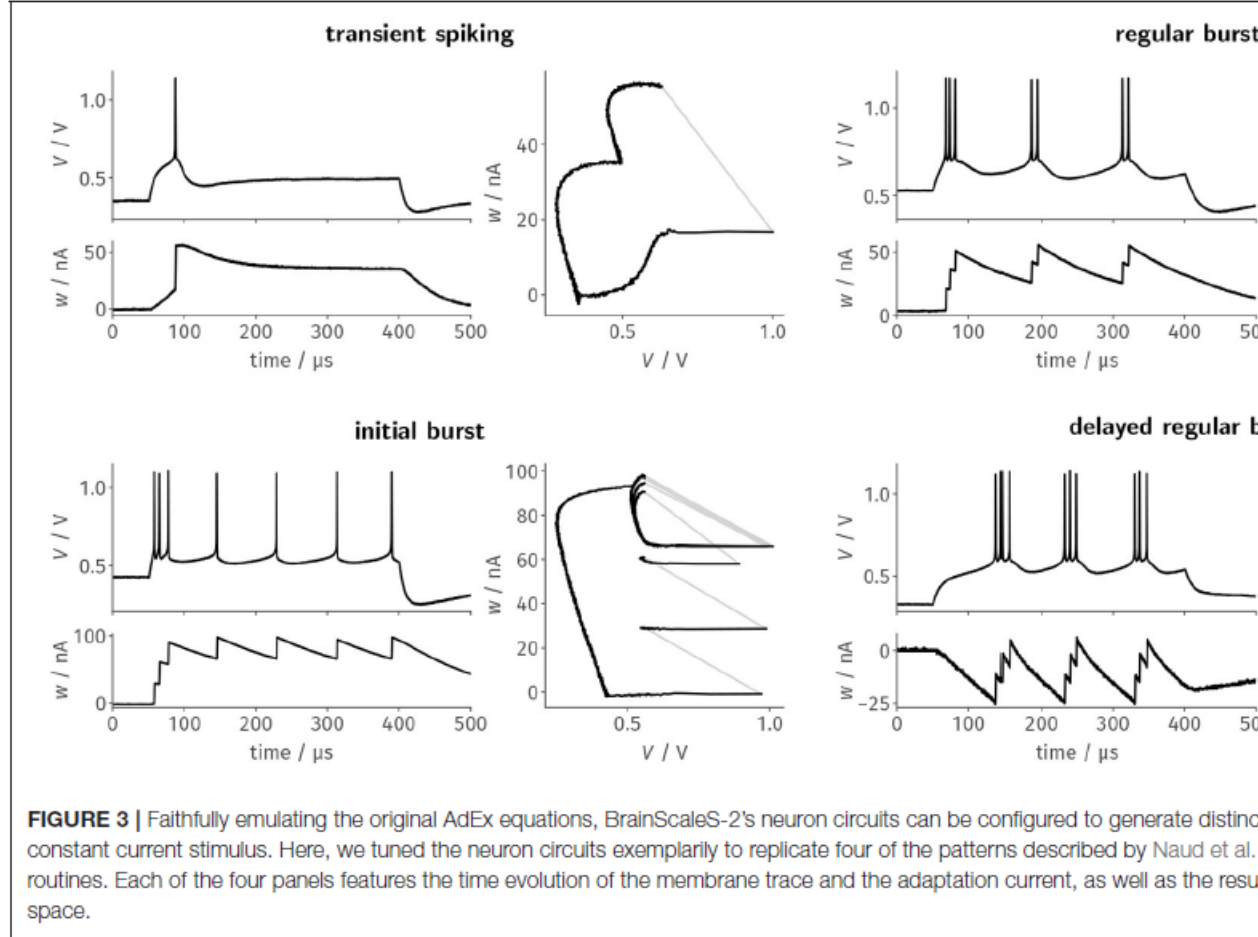
- The plasticity programs can perform both fixed point and integer computations on either a 128 x 8 bit or 64 x 16 bit vector.
- This has been used to demonstrate several versions of STDP.
- Out-of-order-execution is used to speed up the plasticity computations, which allows for vector computations and scalar computations to occur simultaneously.
- The parallel access to the neuron outputs with the column ADC allows for efficient on-chip calibration.
- The 32-bit POWER ISA is used with a compiler and the C++ programming language.
- The SIMD (Single Instruction, Multiple Data) vector instructions are custom but very similar to the POWER ISA ones.
- They modified the GCC compiler to work with their modified ISA and added support for the C++ standard library.
- While on-chip memory is limited to 16 KiB per processor core, they have access to an external memory interface, which is higher latency but larger in size.

Applications of the BrainScaleS-2 System

Faithful Emulation of Complex Dynamics

- The BrainScaleS system allows for detailed control of each circuit and in turn each model parameter.
- The on-chip ADC is 10-bits and operates at a frequency of 29 MHz.
- A FPGA reads and buffers the ADC data, and also handles experiment control.
- For each firing pattern, the neurons were stimulated with a current pulse of width 350 us, and they measured the adaptation voltage (from which current can be estimated) and the membrane potential.

- Four different spiking patterns were observed by varying different parameter such as the adaptation strength a and adaptation increments b .



-
- As shown in the graphs of V vs. w , the neuron spiking is very reproducible and the number of spikes for the same stimulus is almost exactly the same for all 128 neurons.
- The BrainScaleS-2 system also supports multi-compartment models, where the neurons are grouped together into 4 compartments, and the output of one compartment is connected to the input of the next.

- A synaptic input was injected to a different compartment each experiment, and as expected, that input traveled to the compartments following it, but was slightly attenuated with each compartment.
- Plateau potentials are also achieved in compartments by modifying some neuron parameters and providing two inputs (one being delayed) in order to trigger a plateau. A plateau was triggered when the latter compartment spikes before the former, and the time between the spikes at a synapse is inversely proportional to the likelihood of a plateau occurring.

Biology Inspired Learning Approaches

- In order for a learning rule to be biologically plausible, it has to be spatially and temporally local. Spatially local means that neurons can only observe what is directly connected to them. By temporally local, it means that the neuron can make a calculation in a very short time period in response to a stimuli, without seeing the global behavior.
- Multiple reinforcement learning rules were realized on the platform, including playing pong and maze navigation.
- Other biological behavior includes the possibility of synaptic pruning and rewiring and sparse connectivity.
- Insect inspired navigation is also demonstrated. The spiking neural network is configured to represent a bee brain, and the bee hunts for food and returns to its nest. The system runs at 1000x real time. The environment of the bee is simulated using on the on-chip processor in the PPU.
- Since the system is accelerated, tuning hyperparameters takes a very small amount of time.
- The host computer implemented the evolutionary algorithm to allow the bee to learn.
- The system was also used for closed-loop robotics. Since the neuron dynamics are very accelerated, the robotic system had to be on microsecond time scale.